

# Parallel and Distributed Computing

## Assignment 4

102303654

### Exercise 1:

#### 1. Problem Statement

Write an MPI program in a ring topology where Process 0 starts with value 100 and sends it to the next process. Each process adds its own rank to the received value before passing it to the next process. The last process sends the value back to Process 0.

#### 2. Output Obtained

Process 0 starts with value 100  
Process 1 received 100  
Process 1 added rank, now 101  
Process 2 received 101  
Process 2 added rank, now 103  
Process 3 received 103  
Process 3 added rank, now 106  
Process 0 received FINAL value: 106

#### 3. Observations

- The message successfully travelled in a ring
- Each process correctly added its rank ( $0+1+2+3 = 6$ ).
- Final value =  $100 + 6 = 106$ , which matches the expected result.
- Point-to-point communication (MPI\_Send and MPI\_Recv) works reliably with proper use of modulo for next/previous rank.

#### 4. Conclusion

This exercise helped understand cyclic communication and how to handle wrap around in distributed processes using modular arithmetic.

### Exercise 2:

#### 1. Problem Statement

Create an array of numbers 1 to 100 in Process 0. Divide the array equally among all processes using MPI\_Scatter, compute local sums, and combine them using MPI\_Reduce to get the global sum. Also calculate the average.

## **2. Output Obtained**

Process 0 local sum = 325  
Process 1 local sum = 950  
Process 2 local sum = 1575  
Process 3 local sum = 2200  
Global Sum = 5050 (Correct: 5050)

Average = 50.5

## **3. Observations**

- The array was perfectly divided (25 elements per process).
- Local sums added up exactly to 5050, which is the correct sum of first 100 natural numbers.
- MPI\_Scatter and MPI\_Reduce worked efficiently for data distribution and aggregation.
- Bonus average calculation was also correct ( $5050/100 = 50.5$ ).

## **4. Conclusion**

This exercise demonstrated how collective communication primitives can be used to parallelize simple reduction operations like summation, which is very useful in real scientific computing.

# **Exercise 3:**

## **1. Problem Statement**

Each process generates 10 random numbers (0–1000), finds its local maximum and minimum, and then uses MPI\_Reduce with MPI\_MAXLOC and MPI\_MINLOC to find the global maximum and minimum along with the process ID where they occurred.

## **2. Output Obtained**

Process 0 → Local Max: 813 | Local Min: 11  
Process 2 → Local Max: 980 | Local Min: 113  
Process 1 → Local Max: 939 | Local Min: 144  
Process 3 → Local Max: 883 | Local Min: 54  
FINAL RESULT:-  
Global Maximum = 980 (from Process 2)  
Global Minimum = 11 (from Process 0)

### **3. Observations**

- All processes generated different random numbers (due to different seeds).
- MPI\_MAXLOC and MPI\_MINLOC correctly returned both the extreme value and the rank of the process.
- Print order is jumbled because processes run in parallel and print at slightly different times — this is normal behaviour in MPI.
- Global results are accurate.

### **4. Conclusion**

This exercise showed the power of special reduction operations (MAXLOC/MINLOC) to not only find the value but also identify which process produced it.

## **Exercise 4:**

### **1. Problem Statement**

Compute the dot product of two vectors  $A = [1,2,3,4,5,6,7,8]$  and  $B = [8,7,6,5,4,3,2,1]$  in parallel. Scatter both vectors, compute partial dot products in each process, and combine them using MPI\_Reduce.

### **2. Output Obtained**

```
Process 1 partial dot = 38
Process 2 partial dot = 38
Process 0 partial dot = 22
Process 3 partial dot = 22
Final Dot Product = 120 (Correct: 120)
```

### **3. Observations**

- Vectors were correctly divided among 4 processes (2 elements each).
- Partial results:  $22 + 38 + 38 + 22 = 120$  (exact expected value).
- MPI\_Scatter and MPI\_Reduce handled data distribution and summation efficiently.
- Print order was jumbled due to parallel execution.

### **4. Conclusion**

This exercise demonstrated a practical application of parallel matrix/vector operations, which are widely used in machine learning and scientific simulations.